

Lab 2

Brendon Swanepoel - 601949,
Anita de Mello Koch - 1371116,
Nicholas Kastanos - 1393410

1 Comparisons between openmp and Pthreads

$N_0 = N_1$	Basic (s)	Pthread - Diagonal (s)	Pthreads - Blocked (s)	OpenMP - Naive (s)	OpenMP - Diagonal (s)	OpenMP - Blocked (s)
128	5.5828×10^{-5}	0.0057640	0.031893	0.0031893	5.3747×10^{-5}	0.00067019
1024	0.0037808	0.051978	5.2157	0.0026004	0.0019828	0.0022198
2048	0.021078	0.18114	21.217	0.0061313	0.0055345	0.0046060
4096	0.088854	0.63340	85.449	0.025705	0.02354	0.021988
16384	1.7036	9.5218	-	1.2415	1.2145	0.43899

2 Parallel Threading Algorithm

In order to parallelise the transposition algorithm, the transposition is completed one row at a time in a separate thread. To avoid swapping elements which have already been swapped, the algorithm makes use of a nested for loop with a depth of two. The outer loop row counter i runs through every row, and the inner loop counter for the columns begins at $j=i+1$.

2.1 PThreads

Using Pthreads, each column is given to a single thread. This thread swaps that row and column, from $j = i + 1$ until the length of the matrix.

2.2 OpenMP

2.2.1 Naive method

Using the naive method of parallelization, both the inner and outer loops are parallelized. This gives each thread an single swap to complete. This is done by collapsing the for loops.

2.2.2 Diagonal method

Using the diagonal method, only the outer for loop is parallelized. As a result, each thread swaps a single column from the diagonal to the rightmost element.

3 Block Transposition Algorithm

The Block Transposition algorithm completes the process by swapping two blocks and then each block is transposed. This is shown to result in transposition as shown below:

$$\begin{bmatrix} A & B \\ C & D \end{bmatrix}^T = \begin{bmatrix} \begin{bmatrix} A \\ C \end{bmatrix}^T \\ \begin{bmatrix} B \\ D \end{bmatrix}^T \end{bmatrix} = \begin{bmatrix} A^T & C^T \\ B^T & D^T \end{bmatrix}$$

3.1 PThreads

Creating this operation using PThreads, each block B_{ij} of 2×2 matrix elements are assigned to a thread and is swapped with the corresponding block B_{ji} . Once this operation has been completed, two more child threads are spawned to transpose B_{ij} and B_{ji} .

3.2 OpenMP

The parallelisation using OpenMP involves parallelising the block operations. Each thread uses additional parallelisation to swap B_{ij} and B_{ji} and subsequently each transposition is parallelised.

4 Discussion

From the results, for smaller matrix sizes, the non-threaded implementations outperform the threaded implementations, excluding the diagonal implementation for OpenMP. For these smaller matrix sizes, the loss in time due to overhead outweighs the benefit gained from parallelizing the programme. As the matrix size increases, it is clear the parallelized implementations outperform the non-threaded, with exception of the Pthread codes.

For OpenMP, the block transposition method only begins to outperform the naive and diagonal methods when the matrix size increases. This is because the increase in overhead outweighs the benefits gained of block transposition until the matrix is large. The diagonal outperforms the naive implementation due to the decrease in overhead in the diagonal method when compared to the naive method.

OpenMP outperforms Pthreads for each implementation. OpenMP is better optimized than Pthreads for the C++ language, resulting in this difference.

The Pthread block algorithm did not perform very well due to the way in which the threads were generated and how the memory was being allocated. A better way to perform the task is to create a queue of data which is supplied to a set number of threads. Once a thread is no longer actively running, it is assigned a new task until the transposition is completed.

5 Pseudocode

```
input : Pointer to a 2D square matrix
output: In-place transposition of matrix
for Each row of the matrix do
    | for Each column element after the current row value do
    | | Transpose current row and column elements;
    | end
end
```

Algorithm 1: Basic Transposition Algorithm

```
input : Pointer to a 2D square matrix
output: In-place transposition of matrix
Create array of Pthreads of the same size as the matrix dimension;
Create data array of structs for the threads to work on;
Initialise the Pthread attribute to joinable;
Populate data with pointers to the matrix and row numbers;
for Each row of the matrix do
    | Create thread for each row of matrix;
    | Call transposistion function;
end
for Each row of the matrix do
    | Join the created threads;
end
```

Transposition function

```
input : Pointer to thread argument
output: Transposed diagonal
Extract data from thread argument pointer;
Extract current row from thread argument pointer;
for Each column element after the current row value do
    | Transpose along current diagonal;
end
```

Algorithm 2: Diagonal Pthread Transposition Algorithm

input : Pointer to a 2D square matrix, matrix size, block size
output: In-place transposition of matrix
 Create array of Pthreads of the of size $((\text{matrixSize}/\text{blockSize}) * (\text{matrixSize}/\text{blockSize})/2) + ((\text{matrixSize}/\text{blockSize})/2)$;
 Create data array of structs for the threads to work on; Initialise the Pthread attribute to joinable;
 Populate data with pointers to the matrix and block size;
 Initialise independent iterator to 0;
for *Each row of the matrix in steps of the block size* **do**
 | **for** *Each column of the matrix in steps of the block size* **do**
 | | Populate data array with row and column indices;
 | | Create thread for each row of matrix;
 | | Call transposition function;
 | | Increase independent iterator by 1;
 | **end**
end
for *Each row of the matrix* **do**
 | Join the created threads;
end

Transposition function

input : Pointer to thread argument
output: Transposed diagonal
 Extract data from thread argument pointer;
 Extract current row from thread argument pointer;
for *The size of the block* **do**
 | **for** *The size of the block* **do**
 | | Transpose the block elements;
 | **end**
end

Algorithm 3: Block Pthread Transposition Algorithm

input : Pointer to a 2D square matrix
output: In-place transposition of matrix
 OMP Parallel For Loop;
for *Each row of the matrix* **do**
 | OMP Parallel For Loop;
 | **for** *Each column of the matrix* **do**
 | | Transpose matrix emelents;
 | **end**
end

Algorithm 4: Naive OpenMP Transposition Algorithm

input : Pointer to a 2D square matrix
output: In-place transposition of matrix
 OMP Parallel For Loop;
for *Each row of the matrix* **do**
 | **for** *Each column of the matrix* **do**
 | | Transpose matrix emelents;
 | **end**
end

Algorithm 5: Diagonal OpenMP Transposition Algorithm

```

input  : Pointer to a 2D square matrix
output: In-place transposition of matrix
OMP Parallel For Loop;
for Each row of the matrix in steps of the block size do
    for Each column of the matrix in steps of the block size do
        OMP Parallel For Loop;
        for The block size do
            for The block size do
                | Transpose block emelents;
            end
        end
    end
end

```

Algorithm 6: Block OpenMP Transposition Algorithm